

# Convolutional Neural Networks

November 1, 2018

## 1 Lab 11

## 2 Convolutional neural networks

2.1 Submitted to: Prof. Sweetlin Hemlatha

2.2 Submitted by: Prateek Singh (15BCE1091)

```
In [1]: from __future__ import print_function
import numpy as np
import tensorflow as tf
from six.moves import cPickle as pickle
from six.moves import range
```

```
In [2]: url = 'https://commondatastorage.googleapis.com/books1000/'
last_percent_reported = None
data_root = '.' # Change me to store data elsewhere
```

```
def download_progress_hook(count, blockSize, totalSize):
    global last_percent_reported
    percent = int(count * blockSize * 100 / totalSize)
```

```
    if last_percent_reported != percent:
        if percent % 5 == 0:
            sys.stdout.write("%s%%" % percent)
            sys.stdout.flush()
        else:
            sys.stdout.write(".")
            sys.stdout.flush()

    last_percent_reported = percent
```

```
def maybe_download(filename, expected_bytes, force=False):
```

```
    """Download a file if not present, and make sure it's the right size."""
```

```
    dest_filename = os.path.join(data_root, filename)
```

```
    if force or not os.path.exists(dest_filename):
```

```
        print('Attempting to download:', filename)
```

```
        filename, _ = urlretrieve(url + filename, dest_filename, reporthook=download_progr
```

```

    print('\nDownload Complete!')
    statinfo = os.stat(dest_filename)
    if statinfo.st_size == expected_bytes:
        print('Found and verified', dest_filename)
    else:
        raise Exception(
            'Failed to verify ' + dest_filename + '. Can you get to it with a browser?')
    return dest_filename

train_filename = maybe_download('notMNIST_large.tar.gz', 247336696)
test_filename = maybe_download('notMNIST_small.tar.gz', 8458043)

Found and verified ./notMNIST_large.tar.gz
Found and verified ./notMNIST_small.tar.gz

```

Extract the dataset from the compressed .tar.gz file. This would give us a set of directories, labeled A through J.

```

In [3]: num_classes = 10
        np.random.seed(133)

def maybe_extract(filename, force=False):
    root = os.path.splitext(os.path.splitext(filename)[0])[0] # remove .tar.gz
    if os.path.isdir(root) and not force:
        # You may override by setting force=True.
        print('%s already present - Skipping extraction of %s.' % (root, filename))
    else:
        print('Extracting data for %s. This may take a while. Please wait.' % root)
        tar = tarfile.open(filename)
        sys.stdout.flush()
        tar.extractall(data_root)
        tar.close()
    data_folders = [
        os.path.join(root, d) for d in sorted(os.listdir(root))
        if os.path.isdir(os.path.join(root, d))]
    if len(data_folders) != num_classes:
        raise Exception(
            'Expected %d folders, one per class. Found %d instead.' % (
                num_classes, len(data_folders)))
    print(data_folders)
    return data_folders

train_folders = maybe_extract(train_filename)
test_folders = maybe_extract(test_filename)

./notMNIST_large already present - Skipping extraction of ./notMNIST_large.tar.gz.
['./notMNIST_large/A', './notMNIST_large/B', './notMNIST_large/C', './notMNIST_large/D', './notMNIST_large/E', './notMNIST_large/F', './notMNIST_large/G', './notMNIST_large/H', './notMNIST_large/I', './notMNIST_large/J']
./notMNIST_small already present - Skipping extraction of ./notMNIST_small.tar.gz.

```

```
[ './notMNIST_small/A', './notMNIST_small/B', './notMNIST_small/C', './notMNIST_small/D', './notMNIST_small/E', './notMNIST_small/F', './notMNIST_small/G', './notMNIST_small/H', './notMNIST_small/I', './notMNIST_small/J', './notMNIST_small/K', './notMNIST_small/L', './notMNIST_small/M', './notMNIST_small/N', './notMNIST_small/O', './notMNIST_small/P', './notMNIST_small/Q', './notMNIST_small/R', './notMNIST_small/S', './notMNIST_small/T', './notMNIST_small/U', './notMNIST_small/V', './notMNIST_small/W', './notMNIST_small/X', './notMNIST_small/Y', './notMNIST_small/Z' ]
```

Looking at the dataset to ensure it's sensible

```
In [4]: for folders in train_folders:
        a = os.listdir(folders)
        display(Image(filename = folders + '/' + a[0]))
```





Loading the data in a more manageable format

We'll convert the entire dataset into a 3D array (image index, x, y) of floating point values, normalized to have approximately zero mean and standard deviation ~0.5 to make training easier down the road.

A few images might not be readable, we'll just skip them.

```
In [5]: image_size = 28  # Pixel width and height.
        pixel_depth = 255.0  # Number of levels per pixel.

def load_letter(folder, min_num_images):
    """Load the data for a single letter label."""
    image_files = os.listdir(folder)
    dataset = np.ndarray(shape=(len(image_files), image_size, image_size),
                          dtype=np.float32)

    print(folder)
    num_images = 0
    for image in image_files:
        image_file = os.path.join(folder, image)
        try:
            image_data = (imageio.imread(image_file).astype(float) -
                          pixel_depth / 2) / pixel_depth
            if image_data.shape != (image_size, image_size):
                raise Exception('Unexpected image shape: %s' % str(image_data.shape))
            dataset[num_images, :, :] = image_data
            num_images = num_images + 1
        except (IOError, ValueError) as e:
            print('Could not read:', image_file, ':', e, '- it\'s ok, skipping.')

    dataset = dataset[0:num_images, :, :]
    if num_images < min_num_images:
        raise Exception('Many fewer images than expected: %d < %d' %
```

```

        (num_images, min_num_images))

    print('Full dataset tensor:', dataset.shape)
    print('Mean:', np.mean(dataset))
    print('Standard deviation:', np.std(dataset))
    return dataset

def maybe_pickle(data_folders, min_num_images_per_class, force=False):
    dataset_names = []
    for folder in data_folders:
        set_filename = folder + '.pickle'
        dataset_names.append(set_filename)
        if os.path.exists(set_filename) and not force:
            # You may override by setting force=True.
            print('%s already present - Skipping pickling.' % set_filename)
        else:
            print('Pickling %s.' % set_filename)
            dataset = load_letter(folder, min_num_images_per_class)
            try:
                with open(set_filename, 'wb') as f:
                    pickle.dump(dataset, f, pickle.HIGHEST_PROTOCOL)
            except Exception as e:
                print('Unable to save data to', set_filename, ':', e)

    return dataset_names

train_datasets = maybe_pickle(train_folders, 45000)
test_datasets = maybe_pickle(test_folders, 1800)

./notMNIST_large/A.pickle already present - Skipping pickling.
./notMNIST_large/B.pickle already present - Skipping pickling.
./notMNIST_large/C.pickle already present - Skipping pickling.
./notMNIST_large/D.pickle already present - Skipping pickling.
./notMNIST_large/E.pickle already present - Skipping pickling.
./notMNIST_large/F.pickle already present - Skipping pickling.
./notMNIST_large/G.pickle already present - Skipping pickling.
./notMNIST_large/H.pickle already present - Skipping pickling.
./notMNIST_large/I.pickle already present - Skipping pickling.
./notMNIST_large/J.pickle already present - Skipping pickling.
./notMNIST_small/A.pickle already present - Skipping pickling.
./notMNIST_small/B.pickle already present - Skipping pickling.
./notMNIST_small/C.pickle already present - Skipping pickling.
./notMNIST_small/D.pickle already present - Skipping pickling.
./notMNIST_small/E.pickle already present - Skipping pickling.
./notMNIST_small/F.pickle already present - Skipping pickling.
./notMNIST_small/G.pickle already present - Skipping pickling.
./notMNIST_small/H.pickle already present - Skipping pickling.
./notMNIST_small/I.pickle already present - Skipping pickling.

```

./notMNIST\_small/J.pickle already present - Skipping pickling.

Verifying that the data still looks good. Displaying a sample of the labels and images from the ndarray.


```
In [6]: import random
def disp_8_img(imgs, titles):
    """Display subplot with 8 images or less"""
    for i, img in enumerate(imgs):
        plt.subplot(2, 4, i+1)
        plt.title(titles[i])
        plt.axis('off')
        plt.imshow(img)

def disp_sample_pickles(data_folders):
    folder = random.sample(data_folders, 1)
    pickle_filename = ''.join(folder) + '.pickle'
    print(folder)
    print(data_folders)
    try:
        with open(pickle_filename, 'rb') as f:
            dataset = pickle.load(f)
    except Exception as e:
        print('Unable to read data from', pickle_filename, ':', e)
        return
    # display
    plt.suptitle(''.join(folder)[-1])
    for i, img in enumerate(random.sample(list(dataset), 8)):
        plt.subplot(2, 4, i+1)
        plt.axis('off')
        plt.imshow(img)

disp_sample_pickles(train_folders)

['./notMNIST_large/J']
['./notMNIST_large/A', './notMNIST_large/B', './notMNIST_large/C', './notMNIST_large/D', './notMNIST_large/E', './notMNIST_large/F', './notMNIST_large/G', './notMNIST_large/H', './notMNIST_large/I', './notMNIST_large/K', './notMNIST_large/L', './notMNIST_large/M', './notMNIST_large/N', './notMNIST_large/O', './notMNIST_large/P', './notMNIST_large/Q', './notMNIST_large/R', './notMNIST_large/S', './notMNIST_large/T', './notMNIST_large/U', './notMNIST_large/V', './notMNIST_large/W', './notMNIST_large/X', './notMNIST_large/Y', './notMNIST_large/Z']
```

J



checking for the data to be balanced across classes.

```
In [7]: def disp_number_images(data_folders):
        for folder in data_folders:
            pickle_filename = ''.join(folder) + '.pickle'
            try:
                with open(pickle_filename, 'rb') as f:
                    dataset = pickle.load(f)
            except Exception as e:
                print('Unable to read data from', pickle_filename, ': ', e)
                return
            print('Number of images in ', folder, ' : ', len(dataset))
```

```
disp_number_images(train_folders)
disp_number_images(test_folders)
```

```
Number of images in ./notMNIST_large/A : 52909
Number of images in ./notMNIST_large/B : 52911
Number of images in ./notMNIST_large/C : 52912
Number of images in ./notMNIST_large/D : 52911
Number of images in ./notMNIST_large/E : 52912
Number of images in ./notMNIST_large/F : 52912
Number of images in ./notMNIST_large/G : 52912
Number of images in ./notMNIST_large/H : 52912
Number of images in ./notMNIST_large/I : 52912
Number of images in ./notMNIST_large/J : 52911
```

```

Number of images in ./notMNIST_small/A : 1872
Number of images in ./notMNIST_small/B : 1873
Number of images in ./notMNIST_small/C : 1873
Number of images in ./notMNIST_small/D : 1873
Number of images in ./notMNIST_small/E : 1873
Number of images in ./notMNIST_small/F : 1872
Number of images in ./notMNIST_small/G : 1872
Number of images in ./notMNIST_small/H : 1872
Number of images in ./notMNIST_small/I : 1872
Number of images in ./notMNIST_small/J : 1872

```

```

In [8]: def make_arrays(nb_rows, img_size):
        if nb_rows:
            dataset = np.ndarray((nb_rows, img_size, img_size), dtype=np.float32)
            labels = np.ndarray(nb_rows, dtype=np.int32)
        else:
            dataset, labels = None, None
        return dataset, labels

def merge_datasets(pickle_files, train_size, valid_size=0):
    num_classes = len(pickle_files)
    print(pickle_files)
    valid_dataset, valid_labels = make_arrays(valid_size, image_size)
    train_dataset, train_labels = make_arrays(train_size, image_size)
    vsize_per_class = valid_size // num_classes
    tsize_per_class = train_size // num_classes

    start_v, start_t = 0, 0
    end_v, end_t = vsize_per_class, tsize_per_class
    end_l = vsize_per_class + tsize_per_class
    for label, pickle_file in enumerate(pickle_files):
        try:
            with open(pickle_file, 'rb') as f:
                letter_set = pickle.load(f)
                # let's shuffle the letters to have random validation and training set
                np.random.shuffle(letter_set)
                if valid_dataset is not None:
                    valid_letter = letter_set[:vsize_per_class, :, :]
                    valid_dataset[start_v:end_v, :, :] = valid_letter
                    valid_labels[start_v:end_v] = label
                    start_v += vsize_per_class
                    end_v += vsize_per_class

                train_letter = letter_set[vsize_per_class:end_l, :, :]
                train_dataset[start_t:end_t, :, :] = train_letter
                train_labels[start_t:end_t] = label
                start_t += tsize_per_class

```



```

        end_t += tsize_per_class
    except Exception as e:
        print('Unable to process data from', pickle_file, ':', e)
        raise

    return valid_dataset, valid_labels, train_dataset, train_labels

train_size = 200000
valid_size = 10000
test_size = 10000

valid_dataset, valid_labels, train_dataset, train_labels = merge_datasets(
    train_datasets, train_size, valid_size)
_, _, test_dataset, test_labels = merge_datasets(test_datasets, test_size)

print('Training:', train_dataset.shape, train_labels.shape)
print('Validation:', valid_dataset.shape, valid_labels.shape)
print('Testing:', test_dataset.shape, test_labels.shape)

['./notMNIST_large/A.pickle', './notMNIST_large/B.pickle', './notMNIST_large/C.pickle', './notMNIST_large/D.pickle', './notMNIST_large/E.pickle',
 './notMNIST_large/F.pickle', './notMNIST_large/G.pickle', './notMNIST_large/H.pickle', './notMNIST_large/I.pickle', './notMNIST_large/J.pickle',
 './notMNIST_small/A.pickle', './notMNIST_small/B.pickle', './notMNIST_small/C.pickle', './notMNIST_small/D.pickle', './notMNIST_small/E.pickle',
 './notMNIST_small/F.pickle', './notMNIST_small/G.pickle', './notMNIST_small/H.pickle', './notMNIST_small/I.pickle', './notMNIST_small/J.pickle']
Training: (200000, 28, 28) (200000,)
Validation: (10000, 28, 28) (10000,)
Testing: (10000, 28, 28) (10000,)

```

Next, we'll randomize the data. It's important to have the labels well shuffled for the training and test distributions to match.

```

In [9]: def randomize(dataset, labels):
        permutation = np.random.permutation(labels.shape[0])
        shuffled_dataset = dataset[permutation,:,:]
        shuffled_labels = labels[permutation]
        return shuffled_dataset, shuffled_labels
train_dataset, train_labels = randomize(train_dataset, train_labels)
test_dataset, test_labels = randomize(test_dataset, test_labels)
valid_dataset, valid_labels = randomize(valid_dataset, valid_labels)

```

```

In [2]: pickle_file = 'notMNIST.pickle'

with open(pickle_file, 'rb') as f:
    save = pickle.load(f)
    train_dataset = save['train_dataset']
    train_labels = save['train_labels']
    valid_dataset = save['valid_dataset']
    valid_labels = save['valid_labels']
    test_dataset = save['test_dataset']
    test_labels = save['test_labels']

```

```

del save # hint to help gc free up memory
print('Training set', train_dataset.shape, train_labels.shape)
print('Validation set', valid_dataset.shape, valid_labels.shape)
print('Test set', test_dataset.shape, test_labels.shape)

```

```

Training set (200000, 28, 28) (200000,)
Validation set (10000, 28, 28) (10000,)
Test set (10000, 28, 28) (10000,)

```

Reformat into a TensorFlow-friendly shape: - convolutions need the image data formatted as a cube (width by height by #channels) - labels as float 1-hot encodings.

```

In [3]: image_size = 28
        num_labels = 10
        num_channels = 1 # grayscale

import numpy as np

def reformat(dataset, labels):
    dataset = dataset.reshape(
        (-1, image_size, image_size, num_channels)).astype(np.float32)
    labels = (np.arange(num_labels) == labels[:,None]).astype(np.float32)
    return dataset, labels
train_dataset, train_labels = reformat(train_dataset, train_labels)
valid_dataset, valid_labels = reformat(valid_dataset, valid_labels)
test_dataset, test_labels = reformat(test_dataset, test_labels)
print('Training set', train_dataset.shape, train_labels.shape)
print('Validation set', valid_dataset.shape, valid_labels.shape)
print('Test set', test_dataset.shape, test_labels.shape)

```

```

Training set (200000, 28, 28, 1) (200000, 10)
Validation set (10000, 28, 28, 1) (10000, 10)
Test set (10000, 28, 28, 1) (10000, 10)

```

```

In [4]: def accuracy(predictions, labels):
        return (100.0 * np.sum(np.argmax(predictions, 1) == np.argmax(labels, 1))
                / predictions.shape[0])

```

building a small network with two convolutional layers, followed by one fully connected layer. Convolutional networks are more expensive computationally, so we'll limit its depth and number of fully connected nodes.

```

In [6]: batch_size = 16
        patch_size = 5
        depth = 16
        num_hidden = 64

```

```

graph = tf.Graph()

with graph.as_default():

    # Input data.
    tf_train_dataset = tf.placeholder(
        tf.float32, shape=(batch_size, image_size, image_size, num_channels))
    tf_train_labels = tf.placeholder(tf.float32, shape=(batch_size, num_labels))
    tf_valid_dataset = tf.constant(valid_dataset)
    tf_test_dataset = tf.constant(test_dataset)

    # Variables.
    layer1_weights = tf.Variable(tf.truncated_normal(
        [patch_size, patch_size, num_channels, depth], stddev=0.1))
    layer1_biases = tf.Variable(tf.zeros([depth]))
    layer2_weights = tf.Variable(tf.truncated_normal(
        [patch_size, patch_size, depth, depth], stddev=0.1))
    layer2_biases = tf.Variable(tf.constant(1.0, shape=[depth]))
    layer3_weights = tf.Variable(tf.truncated_normal(
        [image_size // 4 * image_size // 4 * depth, num_hidden], stddev=0.1))
    layer3_biases = tf.Variable(tf.constant(1.0, shape=[num_hidden]))
    layer4_weights = tf.Variable(tf.truncated_normal(
        [num_hidden, num_labels], stddev=0.1))
    layer4_biases = tf.Variable(tf.constant(1.0, shape=[num_labels]))

    # Model.
    def model(data):
        conv = tf.nn.conv2d(data, layer1_weights, [1, 2, 2, 1], padding='SAME')
        hidden = tf.nn.relu(conv + layer1_biases)
        conv = tf.nn.conv2d(hidden, layer2_weights, [1, 2, 2, 1], padding='SAME')
        hidden = tf.nn.relu(conv + layer2_biases)
        shape = hidden.get_shape().as_list()
        reshape = tf.reshape(hidden, [shape[0], shape[1] * shape[2] * shape[3]])
        hidden = tf.nn.relu(tf.matmul(reshape, layer3_weights) + layer3_biases)
        return tf.matmul(hidden, layer4_weights) + layer4_biases

    # Training computation.
    logits = model(tf_train_dataset)
    loss = tf.reduce_mean(
        tf.nn.softmax_cross_entropy_with_logits(labels=tf_train_labels, logits=logits))

    # Optimizer.
    optimizer = tf.train.GradientDescentOptimizer(0.05).minimize(loss)

    # Predictions for the training, validation, and test data.
    train_prediction = tf.nn.softmax(logits)
    valid_prediction = tf.nn.softmax(model(tf_valid_dataset))
    test_prediction = tf.nn.softmax(model(tf_test_dataset))

```

```

In [7]: num_steps = 1001

with tf.Session(graph=graph) as session:
    tf.global_variables_initializer().run()
    print('Initialized')
    for step in range(num_steps):
        offset = (step * batch_size) % (train_labels.shape[0] - batch_size)
        batch_data = train_dataset[offset:(offset + batch_size), :, :, :]
        batch_labels = train_labels[offset:(offset + batch_size), :]
        feed_dict = {tf_train_dataset : batch_data, tf_train_labels : batch_labels}
        _, l, predictions = session.run(
            [optimizer, loss, train_prediction], feed_dict=feed_dict)
        if (step % 50 == 0):
            print('Minibatch loss at step %d: %f' % (step, l))
            print('Minibatch accuracy: %.1f%%' % accuracy(predictions, batch_labels))
            print('Validation accuracy: %.1f%%' % accuracy(
                valid_prediction.eval(), valid_labels))
    print('Test accuracy: %.1f%%' % accuracy(test_prediction.eval(), test_labels))

```

Initialized  
 Minibatch loss at step 0: 3.252700  
 Minibatch accuracy: 0.0%  
 Validation accuracy: 9.9%  
 Minibatch loss at step 50: 1.850138  
 Minibatch accuracy: 31.2%  
 Validation accuracy: 48.7%  
 Minibatch loss at step 100: 1.010035  
 Minibatch accuracy: 62.5%  
 Validation accuracy: 64.9%  
 Minibatch loss at step 150: 1.289849  
 Minibatch accuracy: 68.8%  
 Validation accuracy: 72.0%  
 Minibatch loss at step 200: 0.505446  
 Minibatch accuracy: 81.2%  
 Validation accuracy: 75.6%  
 Minibatch loss at step 250: 1.345587  
 Minibatch accuracy: 56.2%  
 Validation accuracy: 69.6%  
 Minibatch loss at step 300: 0.695081  
 Minibatch accuracy: 75.0%  
 Validation accuracy: 77.0%  
 Minibatch loss at step 350: 0.638521  
 Minibatch accuracy: 75.0%  
 Validation accuracy: 78.8%  
 Minibatch loss at step 400: 0.511915  
 Minibatch accuracy: 81.2%  
 Validation accuracy: 78.8%  
 Minibatch loss at step 450: 0.779062

```
Minibatch accuracy: 75.0%
Validation accuracy: 78.7%
Minibatch loss at step 500: 0.602857
Minibatch accuracy: 81.2%
Validation accuracy: 78.8%
Minibatch loss at step 550: 1.071783
Minibatch accuracy: 75.0%
Validation accuracy: 80.2%
Minibatch loss at step 600: 0.966325
Minibatch accuracy: 56.2%
Validation accuracy: 80.5%
Minibatch loss at step 650: 0.564390
Minibatch accuracy: 81.2%
Validation accuracy: 81.5%
Minibatch loss at step 700: 0.549062
Minibatch accuracy: 81.2%
Validation accuracy: 81.0%
Minibatch loss at step 750: 0.473424
Minibatch accuracy: 87.5%
Validation accuracy: 81.4%
Minibatch loss at step 800: 0.421251
Minibatch accuracy: 87.5%
Validation accuracy: 80.3%
Minibatch loss at step 850: 0.261832
Minibatch accuracy: 93.8%
Validation accuracy: 81.6%
Minibatch loss at step 900: 0.716957
Minibatch accuracy: 68.8%
Validation accuracy: 82.2%
Minibatch loss at step 950: 0.280484
Minibatch accuracy: 93.8%
Validation accuracy: 82.5%
Minibatch loss at step 1000: 0.360719
Minibatch accuracy: 93.8%
Validation accuracy: 81.6%
Test accuracy: 88.5%
```

The convolutional model above uses convolutions with stride 2 to reduce the dimensionality. Replacing the strides by a max pooling operation (`nn.max_pool()`) of stride 2 and kernel size 2.

```
In [8]: batch_size = 16
        patch_size = 5
        depth = 16
        num_hidden = 64

        graph = tf.Graph()
```

```

with graph.as_default():

    # Input data.
    tf_train_dataset = tf.placeholder(
        tf.float32, shape=(batch_size, image_size, image_size, num_channels))
    tf_train_labels = tf.placeholder(tf.float32, shape=(batch_size, num_labels))
    tf_valid_dataset = tf.constant(valid_dataset)
    tf_test_dataset = tf.constant(test_dataset)

    # Variables.
    layer1_weights = tf.Variable(tf.truncated_normal(
        [patch_size, patch_size, num_channels, depth], stddev=0.1))
    layer1_biases = tf.Variable(tf.zeros([depth]))
    layer2_weights = tf.Variable(tf.truncated_normal(
        [patch_size, patch_size, depth, depth], stddev=0.1))
    layer2_biases = tf.Variable(tf.constant(1.0, shape=[depth]))
    layer3_weights = tf.Variable(tf.truncated_normal(
        [image_size // 4 * image_size // 4 * depth, num_hidden], stddev=0.1))
    layer3_biases = tf.Variable(tf.constant(1.0, shape=[num_hidden]))
    layer4_weights = tf.Variable(tf.truncated_normal(
        [num_hidden, num_labels], stddev=0.1))
    layer4_biases = tf.Variable(tf.constant(1.0, shape=[num_labels]))

    # Model.
    def model(data):
        conv1 = tf.nn.conv2d(data, layer1_weights, [1, 1, 1, 1], padding='SAME')
        bias1 = tf.nn.relu(conv1 + layer1_biases)
        pool1 = tf.nn.max_pool(bias1, [1, 2, 2, 1], [1, 2, 2, 1], padding='SAME')
        conv2 = tf.nn.conv2d(pool1, layer2_weights, [1, 1, 1, 1], padding='SAME')
        bias2 = tf.nn.relu(conv2 + layer2_biases)
        pool2 = tf.nn.max_pool(bias2, [1, 2, 2, 1], [1, 2, 2, 1], padding='SAME')
        shape = pool2.get_shape().as_list()
        reshape = tf.reshape(pool2, [shape[0], shape[1] * shape[2] * shape[3]])
        hidden = tf.nn.relu(tf.matmul(reshape, layer3_weights) + layer3_biases)
        return tf.matmul(hidden, layer4_weights) + layer4_biases

    # Training computation.
    logits = model(tf_train_dataset)
    loss = tf.reduce_mean(
        tf.nn.softmax_cross_entropy_with_logits(labels=tf_train_labels, logits=logits))

    # Optimizer.
    optimizer = tf.train.GradientDescentOptimizer(0.05).minimize(loss)

    # Predictions for the training, validation, and test data.
    train_prediction = tf.nn.softmax(logits)
    valid_prediction = tf.nn.softmax(model(tf_valid_dataset))
    test_prediction = tf.nn.softmax(model(tf_test_dataset))

```

```
In [9]: num_steps = 1001
```

```
with tf.Session(graph=graph) as session:
    tf.global_variables_initializer().run()
    print('Initialized')
    for step in range(num_steps):
        offset = (step * batch_size) % (train_labels.shape[0] - batch_size)
        batch_data = train_dataset[offset:(offset + batch_size), :, :, :]
        batch_labels = train_labels[offset:(offset + batch_size), :]
        feed_dict = {tf_train_dataset : batch_data, tf_train_labels : batch_labels}
        _, l, predictions = session.run([optimizer, loss, train_prediction], feed_dict=feed_dict)
        if (step % 50 == 0):
            print('Minibatch loss at step %d: %f' % (step, l))
            print('Minibatch accuracy: %.1f%%' % accuracy(predictions, batch_labels))
            print('Validation accuracy: %.1f%%' % accuracy(
                valid_prediction.eval(), valid_labels))
    print('Test accuracy: %.1f%%' % accuracy(test_prediction.eval(), test_labels))
```

Initialized

Minibatch loss at step 0: 3.045946  
Minibatch accuracy: 12.5%  
Validation accuracy: 10.0%  
Minibatch loss at step 50: 1.833991  
Minibatch accuracy: 31.2%  
Validation accuracy: 46.6%  
Minibatch loss at step 100: 1.044222  
Minibatch accuracy: 56.2%  
Validation accuracy: 62.4%  
Minibatch loss at step 150: 1.086824  
Minibatch accuracy: 68.8%  
Validation accuracy: 70.6%  
Minibatch loss at step 200: 0.626718  
Minibatch accuracy: 81.2%  
Validation accuracy: 73.7%  
Minibatch loss at step 250: 1.310024  
Minibatch accuracy: 62.5%  
Validation accuracy: 74.4%  
Minibatch loss at step 300: 0.477789  
Minibatch accuracy: 75.0%  
Validation accuracy: 77.4%  
Minibatch loss at step 350: 0.495264  
Minibatch accuracy: 81.2%  
Validation accuracy: 79.5%  
Minibatch loss at step 400: 0.597861  
Minibatch accuracy: 75.0%  
Validation accuracy: 79.2%  
Minibatch loss at step 450: 0.637514  
Minibatch accuracy: 81.2%

Validation accuracy: 80.6%  
Minibatch loss at step 500: 0.441908  
Minibatch accuracy: 87.5%  
Validation accuracy: 80.2%  
Minibatch loss at step 550: 0.707909  
Minibatch accuracy: 75.0%  
Validation accuracy: 81.6%  
Minibatch loss at step 600: 1.103696  
Minibatch accuracy: 75.0%  
Validation accuracy: 82.0%  
Minibatch loss at step 650: 0.399387  
Minibatch accuracy: 87.5%  
Validation accuracy: 83.0%  
Minibatch loss at step 700: 0.450146  
Minibatch accuracy: 87.5%  
Validation accuracy: 82.2%  
Minibatch loss at step 750: 0.412422  
Minibatch accuracy: 87.5%  
Validation accuracy: 82.6%  
Minibatch loss at step 800: 0.354055  
Minibatch accuracy: 87.5%  
Validation accuracy: 83.0%  
Minibatch loss at step 850: 0.240893  
Minibatch accuracy: 93.8%  
Validation accuracy: 83.3%  
Minibatch loss at step 900: 0.460379  
Minibatch accuracy: 93.8%  
Validation accuracy: 83.8%  
Minibatch loss at step 950: 0.301897  
Minibatch accuracy: 87.5%  
Validation accuracy: 84.6%  
Minibatch loss at step 1000: 0.424297  
Minibatch accuracy: 87.5%  
Validation accuracy: 83.5%  
Test accuracy: 89.8%