

Grid Search and Pipelining

November 1, 2018

1 Pipelining in machine learning

```
In [2]: import numpy as np
import pandas as pd
import seaborn as sns

from sklearn import tree
from sklearn.externals import joblib
from sklearn.decomposition import PCA
from sklearn.pipeline import Pipeline
from sklearn.datasets import load_iris
from sklearn.metrics import accuracy_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.linear_model import SGDClassifier
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score

import matplotlib.pyplot as plt
%matplotlib inline
```

1.1 Loading the dataset into memory

```
In [17]: iris = load_iris()
X_train, X_test, y_train, y_test = train_test_split(iris.data,
                                                    iris.target,
                                                    test_size=0.2,
                                                    random_state=42)
```

1.2 Constructing pipelines for data processing

```
sklearn.pipeline.Pipeline(steps, memory=None)
```

```
In [18]: pipe_lr = Pipeline([('scl', StandardScaler()),
                              ('pca', PCA(n_components=2)),
                              ('clf', LogisticRegression(random_state=42))])
```

```

pipe_svm = Pipeline([('scl', StandardScaler()),
                     ('pca', PCA(n_components=2)),
                     ('clf', svm.SVC(random_state=42))])

pipe_dt = Pipeline([('scl', StandardScaler()),
                    ('pca', PCA(n_components=2)),
                    ('clf', tree.DecisionTreeClassifier(random_state=42))])

# List of pipelines for ease of iteration
pipelines = [pipe_lr, pipe_svm, pipe_dt]

# Dictionary of pipelines and classifier types for ease of reference
pipe_dict = {0: 'Logistic Regression', 1: 'Support Vector Machine', 2: 'Decision Tree

```

1.3 Fitting the data to the pipelines

```

In [19]: # Fit the pipelines
         for pipe in pipelines:
             pipe.fit(X_train, y_train)

         # Compare accuracies
         for idx, val in enumerate(pipelines):
             print('%s pipeline test accuracy: %.3f' % (pipe_dict[idx], val.score(X_test, y_test)))

         # Identify the most accurate model on test data
         best_acc = 0.0
         best_clf = 0
         best_pipe = ''
         for idx, val in enumerate(pipelines):
             if val.score(X_test, y_test) > best_acc:
                 best_acc = val.score(X_test, y_test)
                 best_pipe = val
                 best_clf = idx
         print('Classifier with best accuracy: %s' % pipe_dict[best_clf])

         # Save pipeline to file
         joblib.dump(best_pipe, 'best_pipeline.pkl', compress=1)
         print('Saved %s pipeline to file' % pipe_dict[best_clf])

```

```

Logistic Regression pipeline test accuracy: 0.933
Support Vector Machine pipeline test accuracy: 0.900
Decision Tree pipeline test accuracy: 0.867
Classifier with best accuracy: Logistic Regression
Saved Logistic Regression pipeline to file

```

2 Integrating pipeline with grid search

```
In [20]: from sklearn.grid_search import GridSearchCV
```

2.0.1 List of parameters for the decision tree classifier

```
In [21]: print('Decision tree model hyperparameters:\n', pipe_dt.steps[2][1].get_params())
```

Decision tree model hyperparameters:

```
{'class_weight': None, 'criterion': 'gini', 'max_depth': None, 'max_features': None, 'max_lea
```

2.1 Selected hyperparameters for grid search

- **criterion** - This is the function used to evaluate the quality of the split; we will use both options available in Scikit-learn: Gini impurity and information gain (entropy)
- **min_samples_leaf** - This is the minimum number of samples required for a valid leaf node; we will use the integer range 1 to 5
- **max_depth** - The is the maximum depth of the tree; we will use the integer range 1 to 5
- **min_samples_split** - This is the minimum number of samples required in order to split a non-leaf node; we will use the integer range 1 to 5
- **presort** - This indicates whether or not to presort the data in order to speed up the location of best splits during fitting; this does not have any effect on the resulting model accuracy (only on training times).

```
In [22]: param_range = [1, 2, 3, 4, 5]
```

```
# Set grid search params
grid_params = [{'clf__criterion': ['gini', 'entropy'],
                  'clf__min_samples_leaf': param_range,
                  'clf__max_depth': param_range,
                  'clf__min_samples_split': param_range[1:],
                  'clf__presort': [True, False]}]
```

```
In [23]: # Construct grid search
gs = GridSearchCV(estimator=pipe_dt,
                  param_grid=grid_params,
                  scoring='accuracy',
                  cv=10)
```

```
In [24]: # Fit using grid search
gs.fit(X_train, y_train)

# Best accuracy
print('Best accuracy: %.3f' % gs.best_score_)

# Best params
print('\nBest params:\n', gs.best_params_)
```

Best accuracy: 0.925

Best params:

```
{'clf__criterion': 'gini', 'clf__max_depth': 2, 'clf__min_samples_leaf': 1, 'clf__min_samples
```

3 Managing multiple models, pipelines and grid searches

In [13]: *# Construct some pipelines*

```
pipe_lr = Pipeline([('scl', StandardScaler()),
                    ('clf', LogisticRegression(random_state=42))])

pipe_lr_pca = Pipeline([('scl', StandardScaler()),
                        ('pca', PCA(n_components=2)),
                        ('clf', LogisticRegression(random_state=42))])

pipe_rf = Pipeline([('scl', StandardScaler()),
                    ('clf', RandomForestClassifier(random_state=42))])

pipe_rf_pca = Pipeline([('scl', StandardScaler()),
                        ('pca', PCA(n_components=2)),
                        ('clf', RandomForestClassifier(random_state=42))])

pipe_svm = Pipeline([('scl', StandardScaler()),
                     ('clf', svm.SVC(random_state=42))])

pipe_svm_pca = Pipeline([('scl', StandardScaler()),
                         ('pca', PCA(n_components=2)),
                         ('clf', svm.SVC(random_state=42))])
```

3.1 Setting up the grid search parameters for each classifier

```
In [13]: param_range = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
        param_range_fl = [1.0, 0.5, 0.1]

        grid_params_lr = [{'clf__penalty': ['l1', 'l2'],
                           'clf__C': param_range_fl,
                           'clf__solver': ['liblinear']}]

        grid_params_rf = [{'clf__criterion': ['gini', 'entropy'],
                           'clf__min_samples_leaf': param_range,
                           'clf__max_depth': param_range,
                           'clf__min_samples_split': param_range[1:]}]

        grid_params_svm = [{'clf__kernel': ['linear', 'rbf'],
                           'clf__C': param_range}]
```

3.2 Execution

```
In [15]: jobs = -1
```

```
gs_lr = GridSearchCV(estimator=pipe_lr,
                      param_grid=grid_params_lr,
                      scoring='accuracy',
                      cv=10)

gs_lr_pca = GridSearchCV(estimator=pipe_lr_pca,
                         param_grid=grid_params_lr,
                         scoring='accuracy',
                         cv=10)

gs_rf = GridSearchCV(estimator=pipe_rf,
                     param_grid=grid_params_rf,
                     scoring='accuracy',
                     cv=10,
                     n_jobs=jobs)

gs_rf_pca = GridSearchCV(estimator=pipe_rf_pca,
                         param_grid=grid_params_rf,
                         scoring='accuracy',
                         cv=10,
                         n_jobs=jobs)

gs_svm = GridSearchCV(estimator=pipe_svm,
                      param_grid=grid_params_svm,
                      scoring='accuracy',
                      cv=10,
                      n_jobs=jobs)

gs_svm_pca = GridSearchCV(estimator=pipe_svm_pca,
                          param_grid=grid_params_svm,
                          scoring='accuracy',
                          cv=10,
                          n_jobs=jobs)
```

```
In [16]: grids = [gs_lr, gs_lr_pca, gs_rf, gs_rf_pca, gs_svm, gs_svm_pca]
```

```
grid_dict = {0: 'Logistic Regression', 1: 'Logistic Regression w/PCA',
             2: 'Random Forest', 3: 'Random Forest w/PCA',
             4: 'Support Vector Machine', 5: 'Support Vector Machine w/PCA'}
```

3.3 Fitting the grid search objects

```
In [19]: print('Performing model optimizations...')
         best_acc = 0.0
```

```

best_clf = 0
best_gs = ''
for idx, gs in enumerate(grid_dicts):
    print('\nEstimator: %s' % grid_dicts[idx])
    # Fit grid search
    gs.fit(X_train, y_train)
    # Best params
    print('Best params: %s' % gs.best_params_)
    # Best training data accuracy
    print('Best training accuracy: %.3f' % gs.best_score_)
    # Predict on test data with best params
    y_pred = gs.predict(X_test)
    # Test data accuracy of model with best params
    print('Test set accuracy score for best params: %.3f ' % accuracy_score(y_test, y_pred))
    # Track best (highest test accuracy) model
    if accuracy_score(y_test, y_pred) > best_acc:
        best_acc = accuracy_score(y_test, y_pred)
        best_gs = gs
        best_clf = idx
print('\nClassifier with best test set accuracy: %s' % grid_dicts[best_clf])

# Save best grid search pipeline to file
dump_file = 'best_gs_pipeline.pkl'
joblib.dump(best_gs, dump_file, compress=1)
print('\nSaved %s grid search pipeline to file: %s' % (grid_dicts[best_clf], dump_file))

```

Performing model optimizations...

```

Estimator: Logistic Regression
Best params: {'clf__C': 1.0, 'clf__penalty': 'l1', 'clf__solver': 'liblinear'}
Best training accuracy: 0.917
Test set accuracy score for best params: 0.967

```

```

Estimator: Logistic Regression w/PCA
Best params: {'clf__C': 0.5, 'clf__penalty': 'l1', 'clf__solver': 'liblinear'}
Best training accuracy: 0.858
Test set accuracy score for best params: 0.933

```

```

Estimator: Random Forest
Best params: {'clf__criterion': 'gini', 'clf__max_depth': 3, 'clf__min_samples_leaf': 2, 'clf__min_samples_split': 2}
Best training accuracy: 0.942
Test set accuracy score for best params: 1.000

```

```

Estimator: Random Forest w/PCA
Best params: {'clf__criterion': 'entropy', 'clf__max_depth': 5, 'clf__min_samples_leaf': 1, 'clf__min_samples_split': 2}
Best training accuracy: 0.917
Test set accuracy score for best params: 0.900

```

```
Estimator: Support Vector Machine
Best params: {'clf__C': 3, 'clf__kernel': 'linear'}
Best training accuracy: 0.967
Test set accuracy score for best params: 0.967
```

```
Estimator: Support Vector Machine w/PCA
Best params: {'clf__C': 4, 'clf__kernel': 'rbf'}
Best training accuracy: 0.925
Test set accuracy score for best params: 0.900
```

Classifier with best test set accuracy: Random Forest

Saved Random Forest grid search pipeline to file: best_gs_pipeline.pkl

4 On my own dataset

In this notebook, First I have done some exploration on the data using matplotlib and seaborn. Then, I use different classifier models to predict the quality of the wine.

1. Random Forest Classifier
2. Stochastic Gradient Descent Classifier
3. Support Vector Classifier(SVC)

Then I use cross validation evaluation technique to optimize the model performance.

1. Grid Search CV
2. Cross Validation Score

```
In [6]: #Loading dataset
        wine = pd.read_csv('../Dataset/winequality-red.csv', sep=';')
```

```
In [7]: #Let's check how the data is distributed
        wine.head()
```

```
Out[7]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality
0	9.4	5

1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5


```
In [8]: #Information about the data columns
        wine.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
fixed acidity          1599 non-null float64
volatile acidity       1599 non-null float64
citric acid            1599 non-null float64
residual sugar         1599 non-null float64
chlorides              1599 non-null float64
free sulfur dioxide    1599 non-null float64
total sulfur dioxide   1599 non-null float64
density                1599 non-null float64
pH                    1599 non-null float64
sulphates              1599 non-null float64
alcohol                1599 non-null float64
quality                1599 non-null int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

4.1 Let's do some plotting to know how the data columns are distributed in the dataset


```
In [9]: #Here we see that fixed acidity does not give any specification to classify the quality
        fig = plt.figure(figsize = (10,6))
        sns.barplot(x = 'quality', y = 'fixed acidity', data = wine)
```

```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7135278940>
```


In [10]: *#Here we see that its quite a downing trend in the volatile acidity as we go higher t*
`fig = plt.figure(figsize = (10,6))`
`sns.barplot(x = 'quality', y = 'volatile acidity', data = wine)`

Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x7f71316a3860>




```
In [11]: #Composition of citric acid go higher as we go higher in the quality of the wine
fig = plt.figure(figsize = (10,6))
sns.barplot(x = 'quality', y = 'citric acid', data = wine)
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x7f713166f0b8>
```




```
In [12]: fig = plt.figure(figsize = (10,6))
sns.barplot(x = 'quality', y = 'residual sugar', data = wine)
```

```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x7f71315e5438>
```




```
In [13]: #Composition of chloride also go down as we go higher in the quality of the wine
fig = plt.figure(figsize = (10,6))
sns.barplot(x = 'quality', y = 'chlorides', data = wine)
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7131588a20>
```




```
In [14]: fig = plt.figure(figsize = (10,6))
         sns.barplot(x = 'quality', y = 'free sulfur dioxide', data = wine)
```

```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7131502be0>
```




```
In [15]: fig = plt.figure(figsize = (10,6))
         sns.barplot(x = 'quality', y = 'total sulfur dioxide', data = wine)
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x7f71314db5f8>
```




```
In [16]: #Sulphates level goes higher with the quality of wine
fig = plt.figure(figsize = (10,6))
sns.barplot(x = 'quality', y = 'sulphates', data = wine)
```

```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x7f71313e27f0>
```



```
In [17]: #Alcohol level also goes higher as te quality of wine increases
fig = plt.figure(figsize = (10,6))
sns.barplot(x = 'quality', y = 'alcohol', data = wine)
```

```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x7f713144dac8>
```



4.2 Preprocessing Data for performing Machine learning algorithms

```
In [18]: #Making binary classificaion for the response variable.
#Dividing wine as good and bad by giving the limit for the quality
bins = (2, 6.5, 8)
group_names = ['bad', 'good']
wine['quality'] = pd.cut(wine['quality'], bins = bins, labels = group_names)
```

```
In [19]: #Now lets assign a labels to our quality variable
label_quality = LabelEncoder()
```


```
In [20]: #Bad becomes 0 and good becomes 1
wine['quality'] = label_quality.fit_transform(wine['quality'])
```

```
In [21]: wine['quality'].value_counts()
```

```
Out[21]: 0    1382
         1     217
         Name: quality, dtype: int64
```

```
In [22]: sns.countplot(wine['quality'])
```

```
Out[22]: <matplotlib.axes._subplots.AxesSubplot at 0x7f713131ffd0>
```



```
In [23]: #Now seperate the dataset as response variable and feature variabes
X = wine.drop('quality', axis = 1)
y = wine['quality']
```

```
In [24]: #Train and Test splitting of data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=42)
```

```
In [25]: #Applying Standard scaling to get optimized result
sc = StandardScaler()
```

```
In [26]: X_train = sc.fit_transform(X_train)
X_test = sc.fit_transform(X_test)
```

4.3 Our training and testing data is ready now to perform machine learning algorithm

4.3.1 Random Forest Classifier

```
In [27]: rfc = RandomForestClassifier(n_estimators=200)
rfc.fit(X_train, y_train)
pred_rfc = rfc.predict(X_test)
```

```
In [28]: #Let's see how our model performed
print(classification_report(y_test, pred_rfc))
```

	precision	recall	f1-score	support
0	0.91	0.96	0.93	273
1	0.65	0.43	0.51	47
avg / total	0.87	0.88	0.87	320

Random forest gives the accuracy of 87%

```
In [29]: #Confusion matrix for the random forest classification
print(confusion_matrix(y_test, pred_rfc))
```

```
[[262  11]
 [ 27  20]]
```

4.4 Stochastic Gradient Decent Classifier

```
In [30]: sgd = SGDClassifier(penalty=None)
sgd.fit(X_train, y_train)
pred_sgd = sgd.predict(X_test)
```

```
/home/prateek/anaconda3/envs/dltf/lib/python3.6/site-packages/sklearn/linear_model/stochastic_
"and default tol will be 1e-3." % type(self), FutureWarning)
```

```
In [31]: print(classification_report(y_test, pred_sgd))
```

	precision	recall	f1-score	support
0	0.89	0.93	0.91	273
1	0.47	0.34	0.40	47
avg / total	0.83	0.85	0.84	320

84% accuracy using stochastic gradient descent classifier

```
In [32]: print(confusion_matrix(y_test, pred_sgd))
```

```
[[255  18]
 [ 31  16]]
```


4.5 Support Vector Classifier

```
In [33]: svc = SVC()
        svc.fit(X_train, y_train)
        pred_svc = svc.predict(X_test)
```

```
In [34]: print(classification_report(y_test, pred_svc))
```

	precision	recall	f1-score	support
0	0.88	0.98	0.93	273
1	0.71	0.26	0.37	47
avg / total	0.86	0.88	0.85	320

Support vector classifier gets 86%

4.6 Grid Search CV

Trying out a pipeline with PCA and standard scaler

```
In [45]: #Finding best parameters for our SVC model
```

```
param = {
    'clf__C': [0.1,0.8,0.9,1,1.1,1.2,1.3,1.4],
    'clf__kernel': ['linear', 'rbf'],
    'clf__gamma' : [0.1,0.8,0.9,1,1.1,1.2,1.3,1.4]
}
```

```
pipeline_svc = Pipeline([('scl', StandardScaler()),
                          ('pca', PCA(n_components=3)),
                          ('clf', SVC())])
```

```
grid_svc = GridSearchCV(estimator=pipeline_svc, param_grid=param, scoring='accuracy',
```

```
In [46]: grid_svc.fit(X_train, y_train)
```

```
Out[46]: GridSearchCV(cv=10, error_score='raise',
                      estimator=Pipeline(memory=None,
                      steps=[('scl', StandardScaler(copy=True, with_mean=True, with_std=True)), ('pca',
                      svd_solver='auto', tol=0.0, whiten=False)), ('clf', SVC(C=1.0, cache_size=200, clas
                      decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
                      max_iter=-1, probability=False, random_state=None, shrinking=True,
                      tol=0.001, verbose=False))]),
                      fit_params=None, iid=True, n_jobs=1,
                      param_grid={'clf__C': [0.1, 0.8, 0.9, 1, 1.1, 1.2, 1.3, 1.4], 'clf__kernel': [
                      pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                      scoring='accuracy', verbose=0)
```

```
In [47]: #Best parameters for our svc model
        grid_svc.best_params_
```

```
Out[47]: {'clf__C': 1.1, 'clf__gamma': 0.8, 'clf__kernel': 'rbf'}
```

```
In [48]: #Let's run our SVC again with the best parameters.
        svc = SVC(C = 1.1, gamma = 0.8, kernel='rbf')
        svc.fit(X_train, y_train)
        pred_svc = svc.predict(X_test)
        print(classification_report(y_test, pred_svc))
```

	precision	recall	f1-score	support
0	0.89	0.99	0.94	273
1	0.88	0.32	0.47	47
avg / total	0.89	0.89	0.87	320

```
In [39]: #Finding best parameters for our SVC model
```

```
param = {
    'clf__C': [0.1,0.8,0.9,1,1.1,1.2,1.3,1.4],
    'clf__kernel': ['linear', 'rbf'],
    'clf__gamma' : [0.1,0.8,0.9,1,1.1,1.2,1.3,1.4]
}
```

```
pipeline_svc = Pipeline([('scl', StandardScaler()),
                          ('clf', SVC())])
```

```
grid_svc = GridSearchCV(estimator=pipeline_svc, param_grid=param, scoring='accuracy',
```

```
In [40]: grid_svc.fit(X_train, y_train)
```

```
Out[40]: GridSearchCV(cv=10, error_score='raise',
                      estimator=Pipeline(memory=None,
                      steps=[('scl', StandardScaler(copy=True, with_mean=True, with_std=True)), ('clf'
decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
max_iter=-1, probability=False, random_state=None, shrinking=True,
tol=0.001, verbose=False))]),
                      fit_params=None, iid=True, n_jobs=1,
                      param_grid={'clf__C': [0.1, 0.8, 0.9, 1, 1.1, 1.2, 1.3, 1.4], 'clf__kernel': [
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
scoring='accuracy', verbose=0)
```

```
In [41]: #Best parameters for our svc model
        grid_svc.best_params_
```

```
Out[41]: {'clf__C': 1.2, 'clf__gamma': 0.9, 'clf__kernel': 'rbf'}
```

```
In [42]: #Let's run our SVC again with the best parameters.
svc2 = SVC(C = 1.2, gamma = 0.9, kernel= 'rbf')
svc2.fit(X_train, y_train)
pred_svc2 = svc2.predict(X_test)
print(classification_report(y_test, pred_svc2))
```

	precision	recall	f1-score	support
0	0.90	0.99	0.94	273
1	0.89	0.34	0.49	47
avg / total	0.90	0.90	0.88	320

4.6.1 SVC without PCA improves from 86% to 90% using Grid Search CV hence we would go forward with that

4.7 Cross Validation Score for random forest and SGD

```
In [44]: #Now lets try to do some evaluation for random forest model using cross validation.
rfc_eval = cross_val_score(estimator = rfc, X = X_train, y = y_train, cv = 10)
rfc_eval.mean()
```

```
Out[44]: 0.9093073326771653
```

4.7.1 Random forest accuracy increases from 87% to 91 % using cross validation score