

Machine Learning Lab 4

Logistic Regression

In machine learning, the logistic model is a widely used statistical model that, in its basic form, uses a logistic function to model a binary dependent variable; many more complex extensions exist. Logistic Regression is a classification algorithm. It is used to predict a binary outcome (1 / 0, Yes / No, True / False) given a set of independent variables. To represent binary / categorical outcome, we use dummy variables. It can also be thought of as a special case of linear regression when the outcome variable is categorical, where we are using log of odds as dependent variable. In simple words, it predicts the probability of occurrence of an event by fitting data to a logit function.

The dataset

The dataset used to perform this experiment is the wine quality dataset, it is a combination of data on two types of wine variants, namely red wine and white wine, of the portuguese "Vinho Verde" wine. The dataset contains information on the parameters for fixed acidity, volatile acidity, citric acid, residual sugar, chlorides, free sulfur dioxide, total sulfur dioxide, density, pH, sulphates, alcohol.

Experiment

In this experiment I used the sklearn's logistic regression algorithms to predict the quality of a wine.

Using the pandas library in I loaded the red wine and white wine datasets into the memory from their respective csv files and then merged the two datasets into one single pandas dataframe.

Using the `pandas.DataFrame.describe()` function in pandas I calculated the various statistical measures of each of the columns of the dataset.

For performing the experiment I started with plotting the scatter plot for each of the features in the dataset with every other feature, this helped to find if there were any features which were linearly separable. In the case of my dataset they were not.

Next used random forests to find the importance of features in the dataset and as to how much each feature contributes to the importance. The two most important features are the fixed acidity and the volatile acidity.

Finally, I applied logistic regression to the dataset using both l1 and l2 penalty and the using the saga and newton-cg solver. I was able to achieve an accuracy of 52 and 53 percent respectively. The reason for this low accuracy was the dataset being skewed with just 20 examples in class 3 and more than 2000 examples in class 6 and this skewness in the dataset was the reason for a bad performance of the logistic regression.

The code and plots can be found in the accompanying jupyter notebook.

Logistic Regression

November 1, 2018

1 Lab 4

2 Logistic Regression

2.0.1 Submitted to: Prof. Sweetlin Hemlatha

2.0.2 Submitted by: Prateek Singh (15BCE1091)

```
In [4]: import random
import numpy as np
import pandas as pd
import seaborn as sn
from sklearn import preprocessing
from sklearn.metrics import confusion_matrix
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import precision_score, recall_score

import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap
%matplotlib inline
```

```
In [5]: white_wine = pd.read_csv('../Dataset/winequality-white.csv', sep=';')
sn.set(style='ticks', color_codes=True, font_scale=1)
```

```
In [6]: white_wine.head()
```

```
Out[6]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.0	0.27	0.36	20.7	0.045	
1	6.3	0.30	0.34	1.6	0.049	
2	8.1	0.28	0.40	6.9	0.050	
3	7.2	0.23	0.32	8.5	0.058	
4	7.2	0.23	0.32	8.5	0.058	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	45.0	170.0	1.0010	3.00	0.45	
1	14.0	132.0	0.9940	3.30	0.49	

2	30.0	97.0	0.9951	3.26	0.44
3	47.0	186.0	0.9956	3.19	0.40
4	47.0	186.0	0.9956	3.19	0.40

	alcohol	quality
0	8.8	6
1	9.5	6
2	10.1	6
3	9.9	6
4	9.9	6

Working with white wine dataset first, in between the assignment I realized there was no need for regression analysis. As we already have a dataset, we just need to apply the sklearn's logistic regression with different function.

Dividing the dataset into training and testing classes.

```
In [7]: print('Rows in white wine dataset: ', len(white_wine.axes[0]))
        # print('Rows in red wine dataset: ', len(red_wine.axes[0]))
```

Rows in white wine dataset: 4898

Let's use the white wine dataset and we will split the dataset into training and testing

```
In [8]: # Doing some analysis over the dataset
        white_wine.head()
```

```
Out[8]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.0	0.27	0.36	20.7	0.045	
1	6.3	0.30	0.34	1.6	0.049	
2	8.1	0.28	0.40	6.9	0.050	
3	7.2	0.23	0.32	8.5	0.058	
4	7.2	0.23	0.32	8.5	0.058	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	45.0	170.0	1.0010	3.00	0.45	
1	14.0	132.0	0.9940	3.30	0.49	
2	30.0	97.0	0.9951	3.26	0.44	
3	47.0	186.0	0.9956	3.19	0.40	
4	47.0	186.0	0.9956	3.19	0.40	

	alcohol	quality
0	8.8	6
1	9.5	6
2	10.1	6
3	9.9	6
4	9.9	6

```
In [9]: classes = sorted(white_wine.quality.unique())

        print('Number of data points for class: \n')

        for cls in classes:
            num_samples = len(white_wine.loc[white_wine.quality == cls])
            print(cls, " :", num_samples)
```

Number of data points for class:

```
3  : 20
4  : 163
5  : 1457
6  : 2198
7  : 880
8  : 175
9  : 5
```

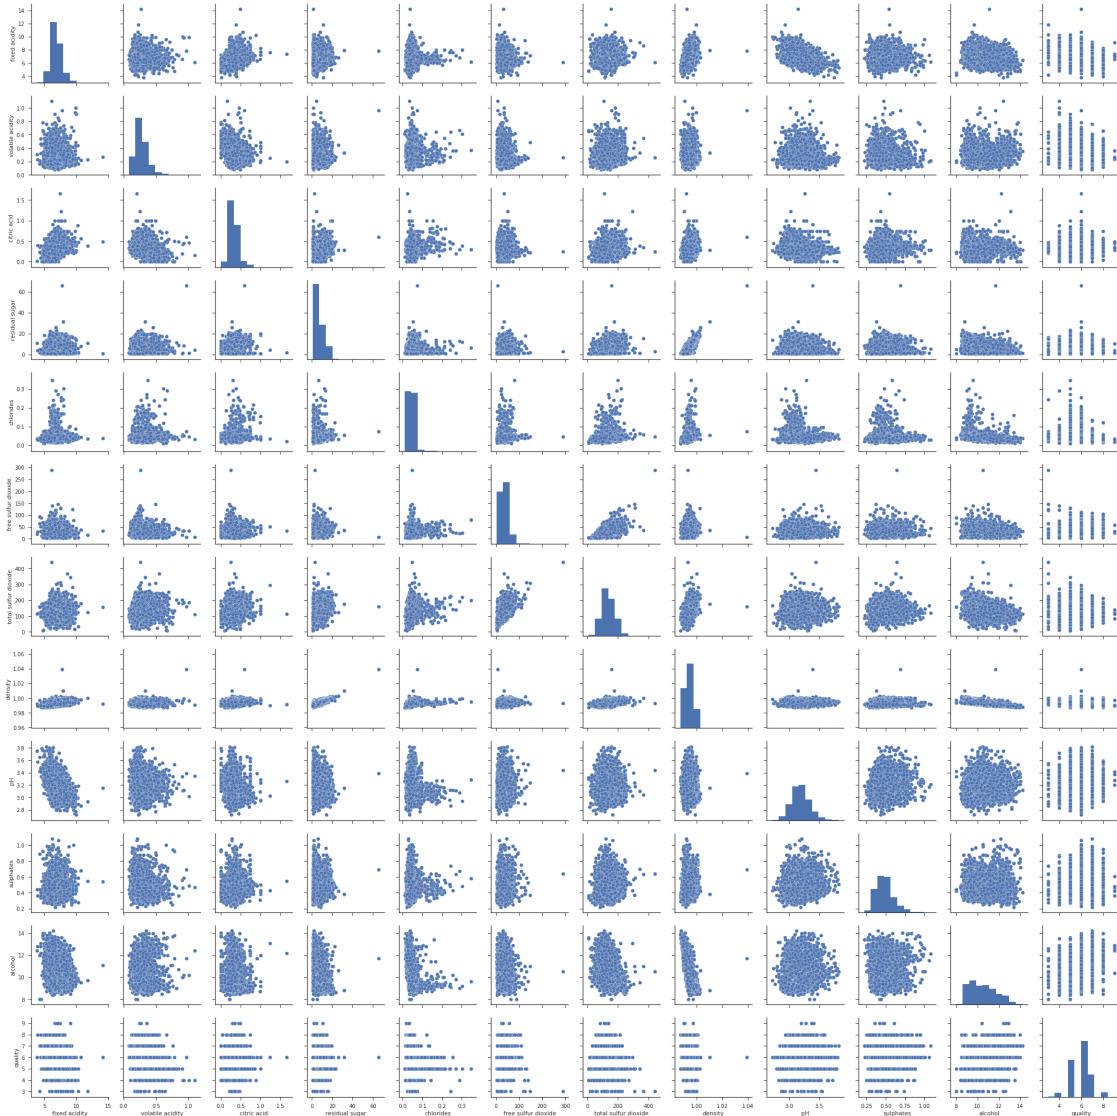
Each data point has 11 features in total, however, we just have around 4400 training example.

Moreover the dataset is very skewed with very few examples in some classes and the major concentration of examples in the other classes.

We either need to reduce the number of features that we're dealing with, by either trying to find the importance of each feature (using decision trees) or dimensionality reduction, or we need to increase the number of examples such that the dataset is less biased towards some of the classes.

```
In [10]: sn.pairplot(white_wine.dropna(), size=2.5)
```

```
Out[10]: <seaborn.axisgrid.PairGrid at 0x7f973e68ca20>
```



```
In [11]: data, labels = white_wine.iloc[:, :11], white_wine.iloc[:, 11]
```

```
forest = ExtraTreesClassifier(n_estimators=250, random_state=0)
forest.fit(data, labels)
importances = forest.feature_importances_
std = np.std([tree.feature_importances_ for tree in forest.estimators_], axis=0)
indices = np.argsort(importances)[::-1]
cumulative_imp = np.cumsum(importances)

print('Feature rankings')

for f in range(data.shape[1]):
    print("%d. feature %d (%f) (%f)" % (f + 1, indices[f], importances[indices[f]], cumulative_imp[indices[f]]))
```

```

plt.figure(figsize=(20, 10))
plt.bar(range(data.shape[1]), importances[indices], color="r", yerr=std[indices], align="center")
plt.xticks(range(data.shape[1]), list(white_wine.columns.values[:11]))
plt.xlim([-1, data.shape[1]])
plt.show()


```

Feature rankings

```

1. feature 10 (0.130143) (1.000000)
2. feature 1 (0.100346) (0.180734)
3. feature 7 (0.092182) (0.701958)
4. feature 5 (0.089962) (0.522192)
5. feature 6 (0.087584) (0.609776)
6. feature 8 (0.085251) (0.787209)
7. feature 3 (0.084598) (0.349213)
8. feature 2 (0.083881) (0.264614)
9. feature 4 (0.083017) (0.432230)
10. feature 9 (0.082648) (0.869857)
11. feature 0 (0.080387) (0.080387)

```



```

In [12]: X_train, X_test, Y_train, Y_test = train_test_split(white_wine.iloc[:, :11],
                                                             white_wine.iloc[:, 11],
                                                             test_size=0.1,
                                                             random_state=42)

print("Size of training set: ", len(X_train.axes[0]))
print("Size of test set: ", len(X_test.axes[0]))

```

```

Size of training set: 4408
Size of test set: 490

```

```
In [13]: X_train.head()
```

```
Out[13]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
1052	7.6	0.29	0.42	1.3	0.035	
3606	6.4	0.38	0.24	7.2	0.047	
1610	7.5	0.32	0.49	1.7	0.031	
621	6.5	0.26	0.43	8.9	0.083	
4750	6.0	0.14	0.37	1.2	0.032	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
1052	18.0	86.0	0.99080	2.99	0.39	
3606	41.0	151.0	0.99604	3.11	0.60	
1610	44.0	109.0	0.99060	3.07	0.46	
621	50.0	171.0	0.99650	2.85	0.50	
4750	63.0	148.0	0.99185	3.32	0.44	

	alcohol
1052	11.3
3606	9.2
1610	12.5
621	9.0
4750	11.2

```
In [14]: # Normalizing the features in Training and testing data.
```

```
X_tr, X_tes = X_train.values, X_test.values
min_max_scaler = preprocessing.MinMaxScaler()
#X_tr_scaled = min_max_scaler.fit_transform(X_tr)
#X_tes_scaled = min_max_scaler.fit_transform(X_tes)
X_train = pd.DataFrame(min_max_scaler.fit_transform(X_tr), index=X_train.index, columns=X_train.columns)
X_test = pd.DataFrame(min_max_scaler.fit_transform(X_tes), index=X_test.index, columns=X_test.columns)

X_train.head()
```

```
Out[14]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
1052	0.4750	0.205882	0.253012	0.010736	0.077151	
3606	0.3250	0.294118	0.144578	0.101227	0.112760	
1610	0.4625	0.235294	0.295181	0.016871	0.065282	
621	0.3375	0.176471	0.259036	0.127301	0.219585	
4750	0.2750	0.058824	0.222892	0.009202	0.068249	

	free sulfur dioxide	total sulfur dioxide	density	pH	\
1052	0.052448	0.178654	0.071139	0.245455	
3606	0.132867	0.329466	0.172161	0.354545	
1610	0.143357	0.232019	0.067284	0.318182	
621	0.164336	0.375870	0.181029	0.118182	
4750	0.209790	0.322506	0.091382	0.545455	

	sulphates	alcohol
1052		
3606		
1610		
621		
4750		

1052	0.197674	0.532258
3606	0.441860	0.193548
1610	0.279070	0.725806
621	0.325581	0.161290
4750	0.255814	0.516129

```
In [15]: forest = ExtraTreesClassifier(n_estimators=250, random_state=0)
forest.fit(X_test, Y_test)
importances = forest.feature_importances_
std = np.std([tree.feature_importances_ for tree in forest.estimators_], axis=0)
cum_imp = np.cumsum(importances)
indices = np.argsort(importances)[::-1]


print('Feature rankings')

for f in range(data.shape[1]):
    print("%d. feature %d (%f)" % (f + 1, indices[f], importances[indices[f]]))

plt.figure(figsize=(20, 10))
plt.bar(range(data.shape[1]), importances[indices], color="r", yerr=std[indices], align="center")
plt.step(range(data.shape[1]), cum_imp, 'b')
plt.xticks(range(data.shape[1]), list(white_wine.columns.values[:11]))
plt.xlim([-1, data.shape[1]])
plt.show()
```

Feature rankings

```
1. feature 10 (0.123645)
2. feature 1 (0.100660)
3. feature 7 (0.097706)
4. feature 5 (0.090431)
5. feature 4 (0.090032)
6. feature 6 (0.086664)
7. feature 8 (0.086057)
8. feature 9 (0.083790)
9. feature 0 (0.080790)
10. feature 2 (0.080779)
11. feature 3 (0.079445)
```



From the above graph it is evident that each of the feature contributes almost equally to the information represented by the dataset and thus omitting any of the features from the logistic regression isn't a wise choice, Hence we will keep all the features in the dataset with us.

```
In [16]: model_l1 = LogisticRegression(penalty='l1', solver='saga', max_iter=100000, multi_class='multinomial')
model_l1.fit(X_train, Y_train)

print("Train Score: ", model_l1.score(X_train, Y_train))
print("Test Score: ", model_l1.score(X_test, Y_test), '\n')

Y_pred = model_l1.predict(X_test)
cm = confusion_matrix(Y_test, Y_pred)
print(cm)
```

Train Score: 0.5410617059891107

Test Score: 0.49183673469387756

```
[[ 0  0  0  3  0  0]
 [ 0  0 11  7  0  0]
 [ 0  0 34 108  1  1]
 [ 0  0 14 188 13  0]
 [ 0  0  3  72 18  1]
 [ 0  0  0  11  4  1]]
```

```
In [17]: model_l2 = LogisticRegression(penalty='l2', solver='newton-cg', max_iter=10000, multi_class='multinomial')
model_l2.fit(X_train, Y_train)

print("Train Score: ", model_l2.score(X_train, Y_train))
print("Test Score: ", model_l2.score(X_test, Y_test), '\n')
```

```

Y_pred = model_l2.predict(X_test)
cm = confusion_matrix(Y_test, Y_pred)
print(cm)

```

Train Score: 0.5394736842105263

Test Score: 0.4959183673469388

```

[[ 0  0  0  3  0  0]
 [ 0  0 10  8  0  0]
 [ 0  0 31 113  0  0]
 [ 0  0 13 193  9  0]
 [ 0  0  2  73 19  0]
 [ 0  0  0  11  5  0]]

```

```

In [18]: def plot_decision_surface(X, y, classifier, test_idx=None, resolution=0.02):

    markers = ('s', 'x', 'o', '^', 'v', '+', '.')
    colors = ('red', 'blue', 'lightgreen', 'gray', 'cyan', 'lightblue', 'lightgreen')
    cmap = ListedColormap(colors[:len(np.unique(y))])

    x1_min, x1_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    x2_min, x2_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, resolution), np.arange(x2_min, x2_max, resolution))

    plt.figure(figsize=(15, 15))
    Z = classifier.predict(np.array([xx1.ravel(), xx2.ravel()]).T)
    Z = Z.reshape(xx1.shape)

    plt.contourf(xx1, xx2, Z, alpha=0.4, cmap=cmap)
    plt.xlim(xx1.min(), xx1.max())
    plt.ylim(xx2.min(), xx2.max())
    plt.xlabel('fixed acidity')
    plt.ylabel('volatile acidity')

    X_test, y_test = X[test_idx, :], y[test_idx]
    for idx, cl in enumerate(np.unique(y)):
        plt.scatter(x=X[y == cl, 0], y=X[y == cl, 1],
                    alpha=0.8, c=cmap(idx),
                    marker=markers[idx], label=cl)
    if test_idx:
        X_test, y_test = X[test_idx, :], y[test_idx]
        plt.scatter(X_test[:, 0], X_test[:, 1], c='',
                    alpha=1.0, linewidth=1, marker='o',
                    s=55, label='test set')

In [19]: model_l2 = LogisticRegression(penalty='l2', solver='newton-cg', max_iter=10000, multi
model_l2.fit(X_train.iloc[:, [1, 10]], Y_train)

```

```

print("Train Score: ", model_l2.score(X_train.iloc[:, [1, 10]], Y_train))
print("Test Score: ", model_l2.score(X_test.iloc[:, [1, 10]], Y_test), '\n')

Y_pred = model_l2.predict(X_test.iloc[:, [1, 10]])
cm = confusion_matrix(Y_test, Y_pred)
print(cm)

```

Train Score: 0.5249546279491834

Test Score: 0.5061224489795918

```

[[ 0  0  2  1  0  0]
 [ 0  1 12  5  0  0]
 [ 0  0 90 53  1  0]
 [ 0  0 64 141 10  0]
 [ 0  0 15 63 16  0]
 [ 0  0  1 11  4  0]]

```

```

In [20]: a = X_train.iloc[:, [1, 10]].values
         plot_decision_surface(X=a, y = np.array(Y_train.values), classifier=model_l2)

```

