

Machine Learning Lab 7

Support Vector Machines

Decision A Support Vector Machine (SVM) is a discriminative classifier formally defined by a separating hyperplane. In other words, given labeled training data (supervised learning), the algorithm outputs an optimal hyperplane which categorizes new examples. In two dimensional space this hyperplane is a line dividing a plane in two parts where in each class lay in either side. Basically support vector machine works on the principle of separation of classes. What SVM does is, it finds out a line/hyper-parameter(in multidimensional space) that separates out classes.

The dataset

The dataset used to perform this experiment is the wine quality dataset, it is a combination of data on two types of wine variants, namely red wine and white wine, of the portuguese "Vinho Verde" wine. The dataset contains information on the parameters for fixed acidity, volatile acidity, citric acid, residual sugar, chlorides, free sulfur dioxide, total sulfur dioxide, density, pH, sulphates, alcohol.

Experiment

In this experiment I used the sklearn's support vector classifier (SVC) algorithm to predict the quality of a wine.

Using the pandas library in I loaded the red wine and white wine datasets into the memory from their respective csv files and then merged the two datasets into one single pandas dataframe.

Using the `pandas.DataFrame.describe()` function in pandas I calculated the various statistical measures of each of the columns of the dataset.

For performing the experiment I started with plotting the scatter plot for each of the features in the dataset with every other feature, this helped to find if there were any features which were linearly separable. In the case of my dataset they were not.

Next I divided the dataset into training and testing portions using the `train_test_split` functionality in sklearn

I setup the grid search for SVC with RBF(Radial Basis function) kernel and the value of gamma ranging from $1e-4$ to 0.5, and the value of parameter C ranging from 1 to 1000. I ran the grid search and for the parameters of $C=1$, and $\gamma=0.5$, I achieve the best performance with a precision score of 0.84.

Next for visualization purposes I selected two most important features from the dataset. In this case the feature importance was decided based on the feature importance values calculated in the previous lab assignments using random forests.

Thus the two features of fixed acidity and volatile acidity were selected and an SVC was trained on this model and the decision boundaries were predicted on the same.

The code and plots can be found in the accompanying jupyter notebook.

Support Vector Machines

November 1, 2018

1 Lab 7

2 Support Vector Machines

2.1 Submitted to: Prof. Sweetlin Hemlatha

2.2 Submitted by: Prateek Singh (15BCE1091)


```
In [16]: import numpy as np
import pandas as pd
import seaborn as sns
from scipy import stats
from sklearn.svm import SVC
from sklearn.utils import shuffle
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.datasets.samples_generator import make_blobs

import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap
%matplotlib inline

sns.set(style='whitegrid', context='notebook', font_scale=1)

In [2]: X, y = make_blobs(n_samples=100, centers=2, random_state=0, cluster_std=0.70)
plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')


Out[2]: <matplotlib.collections.PathCollection at 0x7fd48abf7e10>
```



```
In [3]: xfit = np.linspace(-1, 3.5)
plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')
plt.plot([0.6], [2.1], 'x', color='red', markeredgewidth=2, markersize=10)

for m, b in [(1, 0.65), (0.5, 1.6), (-0.2, 2.9)]:
    plt.plot(xfit, m*xfit + b, '-k')


plt.xlim(-1, 3.5);
```



```
In [4]: xfit = np.linspace(-1, 3.5)
plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')

for m, b, d in [(1, 0.65, 0.33), (0.5, 1.6, 0.55), (-0.2, 2.9, 0.2)]:
    yfit = m*xfit + b
    plt.plot(xfit, yfit, '-k')
    plt.fill_between(xfit, yfit-d, yfit+d, edgecolor='none', color='#AAAAAA', alpha=0.4)

plt.xlim(-1, 3.5);
```



```
In [5]: from sklearn.svm import SVC
        model = SVC(kernel='linear', C=1E10)
        model.fit(X, y)
```

```
Out[5]: SVC(C=10000000000.0, cache_size=200, class_weight=None, coef0=0.0,
          decision_function_shape='ovr', degree=3, gamma='auto', kernel='linear',
          max_iter=-1, probability=False, random_state=None, shrinking=True,
          tol=0.001, verbose=False)
```

```
In [6]: def plot_svc_decision_function(model, ax=None, plot_support=True):
        """Plot the decision function for a 2D SVC"""
        if ax is None:
            ax = plt.gca()
        xlim = ax.get_xlim()
        ylim = ax.get_ylim()

        #create grid to evaluate
        x = np.linspace(xlim[0], xlim[1], 30)
        y = np.linspace(ylim[0], ylim[1], 30)
        Y, X = np.meshgrid(y, x)
        xy = np.vstack([X.ravel(), Y.ravel()]).T
        P = model.decision_function(xy).reshape(X.shape)

        #plot decision boundaries and margins
        ax.contour(X, Y, P, colors='k', levels=[-1, 0, 1], alpha=0.5, linestyle=['--', '-'])
```

```


#plot support vector machines
if plot_support:
    ax.scatter(model.support_vectors_[0], model.support_vectors_[1], s=300, 1.
    ax.set_xlim(xlim)
    ax.set_ylim(ylim)

```

```

In [7]: plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')
        plot_svc_decision_function(model);

```




```

In [8]: from sklearn.datasets.samples_generator import make_circles
        X, y = make_circles(100, factor=.1, noise=.1)


        clf = SVC(kernel='linear').fit(X, y)

        plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')
        plot_svc_decision_function(clf, plot_support=False);

```



```
In [9]: clf = SVC(kernel='rbf', C=1E6)
        clf.fit(X, y)
        plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')
        plot_svc_decision_function(clf, plot_support=False);
```



2.2.1 On my own dataset

```
In [3]: wine_data = pd.read_csv('../Dataset/winequality-white.csv', sep=';')
        red_wine = pd.read_csv('../Dataset/winequality-red.csv', sep=';')
        wine_data.append(red_wine)
        wine_data.head(10)
```

```
Out[3]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.0	0.27	0.36	20.7	0.045	
1	6.3	0.30	0.34	1.6	0.049	
2	8.1	0.28	0.40	6.9	0.050	
3	7.2	0.23	0.32	8.5	0.058	
4	7.2	0.23	0.32	8.5	0.058	
5	8.1	0.28	0.40	6.9	0.050	
6	6.2	0.32	0.16	7.0	0.045	
7	7.0	0.27	0.36	20.7	0.045	
8	6.3	0.30	0.34	1.6	0.049	
9	8.1	0.22	0.43	1.5	0.044	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	45.0	170.0	1.0010	3.00	0.45	
1	14.0	132.0	0.9940	3.30	0.49	
2	30.0	97.0	0.9951	3.26	0.44	
3	47.0	186.0	0.9956	3.19	0.40	
4	47.0	186.0	0.9956	3.19	0.40	
5	30.0	97.0	0.9951	3.26	0.44	
6	30.0	136.0	0.9949	3.18	0.47	
7	45.0	170.0	1.0010	3.00	0.45	
8	14.0	132.0	0.9940	3.30	0.49	
9	28.0	129.0	0.9938	3.22	0.45	

	alcohol	quality
0	8.8	6
1	9.5	6
2	10.1	6
3	9.9	6
4	9.9	6
5	10.1	6
6	9.6	6
7	8.8	6
8	9.5	6
9	11.0	6

```
In [4]: wine_data.dtypes
```

```
Out[4]: fixed acidity      float64
        volatile acidity   float64
        citric acid        float64
        residual sugar     float64
```


```

chlorides          float64
free sulfur dioxide float64
total sulfur dioxide float64
density            float64
pH                float64
sulphates          float64
alcohol            float64
quality            int64
dtype: object

```

```
In [5]: sns.pairplot(wine_data.iloc[:, :11].dropna(), size=2.5)
```

```
Out[5]: <seaborn.axisgrid.PairGrid at 0x7fd48ac0d550>
```



```
In [6]: X_train, X_test, Y_train, Y_test = train_test_split(wine_data.iloc[:, :11],
                                                            wine_data.iloc[:, 11],
                                                            test_size=0.2,
                                                            random_state=42)

print('Size of training set: ', len(X_train.axes[0]))
print('Size of test set: ', len(X_test.axes[0]))
```

```
Size of training set: 3918
Size of test set: 980
```

```
In [7]: parameters = [{'kernel': ['rbf'],
                        'gamma': [1e-4, 1e-3, 0.01, 0.1, 0.2, 0.5],
                        'C': [1, 10, 100, 1000]},
                      {'kernel': ['linear'], 'C': [1, 10, 100, 1000]}]
```

```
In [8]: clf = GridSearchCV(SVC(decision_function_shape='ovr'), parameters, cv=5)
        clf.fit(X_train, Y_train)
```

```
Out[8]: GridSearchCV(cv=5, error_score='raise',
                    estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
                                decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
                                max_iter=-1, probability=False, random_state=None, shrinking=True,
                                tol=0.001, verbose=False),
                    fit_params=None, iid=True, n_jobs=1,
                    param_grid=[{'kernel': ['rbf'], 'gamma': [0.0001, 0.001, 0.01, 0.1, 0.2, 0.5],
                                pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                                scoring=None, verbose=0}]
```

```
In [9]: print("Best parameters set found on development set:")
        print(clf.best_params_)
        print("Grid scores on training set:")
        means = clf.cv_results_['mean_test_score']
        stds = clf.cv_results_['std_test_score']
        for mean, std, params in zip(means, stds, clf.cv_results_['params']):
            print("%0.3f (+/-%0.03f) for %r"
                  % (mean, std * 2, params))
```

```
Best parameters set found on development set:
```

```
{'C': 1, 'gamma': 0.5, 'kernel': 'rbf'}
```

```
Grid scores on training set:
```

```
0.450 (+/-0.016) for {'C': 1, 'gamma': 0.0001, 'kernel': 'rbf'}
0.472 (+/-0.018) for {'C': 1, 'gamma': 0.001, 'kernel': 'rbf'}
0.490 (+/-0.025) for {'C': 1, 'gamma': 0.01, 'kernel': 'rbf'}
0.555 (+/-0.008) for {'C': 1, 'gamma': 0.1, 'kernel': 'rbf'}
0.576 (+/-0.027) for {'C': 1, 'gamma': 0.2, 'kernel': 'rbf'}
0.593 (+/-0.019) for {'C': 1, 'gamma': 0.5, 'kernel': 'rbf'}
0.479 (+/-0.035) for {'C': 10, 'gamma': 0.0001, 'kernel': 'rbf'}
0.503 (+/-0.022) for {'C': 10, 'gamma': 0.001, 'kernel': 'rbf'}
```

```

0.537 (+/-0.016) for {'C': 10, 'gamma': 0.01, 'kernel': 'rbf'}
0.556 (+/-0.018) for {'C': 10, 'gamma': 0.1, 'kernel': 'rbf'}
0.573 (+/-0.044) for {'C': 10, 'gamma': 0.2, 'kernel': 'rbf'}
0.589 (+/-0.029) for {'C': 10, 'gamma': 0.5, 'kernel': 'rbf'}
0.517 (+/-0.019) for {'C': 100, 'gamma': 0.0001, 'kernel': 'rbf'}
0.526 (+/-0.027) for {'C': 100, 'gamma': 0.001, 'kernel': 'rbf'}
0.554 (+/-0.006) for {'C': 100, 'gamma': 0.01, 'kernel': 'rbf'}
0.555 (+/-0.021) for {'C': 100, 'gamma': 0.1, 'kernel': 'rbf'}
0.573 (+/-0.044) for {'C': 100, 'gamma': 0.2, 'kernel': 'rbf'}
0.589 (+/-0.029) for {'C': 100, 'gamma': 0.5, 'kernel': 'rbf'}
0.528 (+/-0.018) for {'C': 1000, 'gamma': 0.0001, 'kernel': 'rbf'}
0.551 (+/-0.029) for {'C': 1000, 'gamma': 0.001, 'kernel': 'rbf'}
0.566 (+/-0.036) for {'C': 1000, 'gamma': 0.01, 'kernel': 'rbf'}
0.556 (+/-0.021) for {'C': 1000, 'gamma': 0.1, 'kernel': 'rbf'}
0.573 (+/-0.044) for {'C': 1000, 'gamma': 0.2, 'kernel': 'rbf'}
0.589 (+/-0.029) for {'C': 1000, 'gamma': 0.5, 'kernel': 'rbf'}
0.526 (+/-0.015) for {'C': 1, 'kernel': 'linear'}
0.526 (+/-0.018) for {'C': 10, 'kernel': 'linear'}
0.534 (+/-0.024) for {'C': 100, 'kernel': 'linear'}
0.522 (+/-0.033) for {'C': 1000, 'kernel': 'linear'}

```

```

In [10]: print("Detailed classification report:")
         print("The model is trained on the full development set.")
         print("The scores are computed on the full evaluation set.")
         #data_train, data_test, label_train, label_test

         y_true, y_pred = Y_test, clf.predict(X_test)
         print(classification_report(y_true, y_pred))
         print()

```

Detailed classification report:

The model is trained on the full development set.

The scores are computed on the full evaluation set.

	precision	recall	f1-score	support
3	0.00	0.00	0.00	5
4	1.00	0.08	0.15	25
5	0.83	0.39	0.53	291
6	0.55	0.95	0.70	432
7	0.91	0.39	0.55	192
8	1.00	0.37	0.54	35
avg / total	0.73	0.63	0.60	980

/home/prateek/anaconda3/envs/dltf/lib/python3.6/site-packages/sklearn/metrics/classification.py

```
'precision', 'predicted', average, warn_for)
```

```
In [21]: def plot_decision_surface(X, y, classifier, test_idx=None, resolution=0.02):

    markers = ('s', 'x', 'o', '^', 'v', '+', '.')
    colors = ('red', 'blue', 'lightgreen', 'gray', 'cyan', 'lightblue', 'lightgreen')
    cmap = ListedColormap(colors[:len(np.unique(y))])

    x1_min, x1_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    x2_min, x2_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, resolution), np.arange(x2_min, x2_max, resolution))

    Z = classifier.predict(np.array([xx1.ravel(), xx2.ravel()]).T)
    Z = Z.reshape(xx1.shape)

    plt.figure(figsize=(15,15))
    plt.contourf(xx1, xx2, Z, alpha=0.4, cmap=cmap)
    plt.xlim(xx1.min(), xx1.max())
    plt.ylim(xx2.min(), xx2.max())
    plt.xlabel('fixed acidity')
    plt.ylabel('volatile acidity')

    X_test, y_test = X[test_idx, :], y[test_idx]

    for idx, cl in enumerate(np.unique(y)):
        plt.scatter(x=X[y == cl, 0], y=X[y == cl, 1],
                    alpha=0.8, c=cmap(idx),
                    marker=markers[idx], label=cl)
        if test_idx:
            X_test, y_test = X[test_idx, :], y[test_idx]
            plt.scatter(X_test[:, 0], X_test[:, 1], c='',
                        alpha=1.0, linewidth=1, marker='o',
                        s=55, label='test set')

In [22]: svc = SVC(C=1.0, kernel='rbf')
        svc.fit(X_train.iloc[:, [1, 10]], Y_train)

Out [22]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
              decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
              max_iter=-1, probability=False, random_state=None, shrinking=True,
              tol=0.001, verbose=False)

In [23]: a = np.array(X_train.iloc[:, [1, 10]])
        plot_decision_surface(X=a, y = np.array(Y_train.values), classifier=svc)
```

