

KMeans + KNN

November 1, 2018

1 Lab 8

2 K-Means Clustering and KNN

2.1 Submitted to: Prof. Sweetlin Hemlatha

2.2 Submitted by: Prateek Singh (15BCE1091)

```
In [1]: import random
import numpy as np
import pandas as pd
import seaborn as sn

from sklearn.cluster import KMeans
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import precision_score, recall_score
from sklearn.neighbors import KNeighborsClassifier

import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap
%matplotlib inline

In [2]: red_wine_data = pd.read_csv('../Dataset/winequality-red.csv', sep=';')
white_wine_data = pd.read_csv('../Dataset/winequality-white.csv', sep=';')

wine_data = pd.concat([red_wine_data, white_wine_data])
bins = (2, 6.5, 10)
group_names = ['bad', 'good']
wine_data['quality'] = pd.cut(wine_data['quality'], bins = bins, labels = group_names)
wine_data.iloc[:, :11].describe()
```

Out [2]:

	fixed acidity	volatile acidity	citric acid	residual sugar	\
count	6497.000000	6497.000000	6497.000000	6497.000000	
mean	7.215307	0.339666	0.318633	5.443235	
std	1.296434	0.164636	0.145318	4.757804	
min	3.800000	0.080000	0.000000	0.600000	

25%	6.400000	0.230000	0.250000	1.800000
50%	7.000000	0.290000	0.310000	3.000000
75%	7.700000	0.400000	0.390000	8.100000
max	15.900000	1.580000	1.660000	65.800000

	chlorides	free sulfur dioxide	total sulfur dioxide	density \
count	6497.000000	6497.000000	6497.000000	6497.000000
mean	0.056034	30.525319	115.744574	0.994697
std	0.035034	17.749400	56.521855	0.002999
min	0.009000	1.000000	6.000000	0.987110
25%	0.038000	17.000000	77.000000	0.992340
50%	0.047000	29.000000	118.000000	0.994890
75%	0.065000	41.000000	156.000000	0.996990
max	0.611000	289.000000	440.000000	1.038980

	pH	sulphates	alcohol
count	6497.000000	6497.000000	6497.000000
mean	3.218501	0.531268	10.491801
std	0.160787	0.148806	1.192712
min	2.720000	0.220000	8.000000
25%	3.110000	0.430000	9.500000
50%	3.210000	0.510000	10.300000
75%	3.320000	0.600000	11.300000
max	4.010000	2.000000	14.900000

```
In [3]: scaler = StandardScaler()
```

```
data = wine_data.iloc[:,11].values
scaled_features = scaler.fit_transform(data)
```

```
wine_data_scaled = pd.DataFrame(scaled_features, index=wine_data.index, columns=wine_data.columns)
wine_data_scaled.head()
```

```
Out [3]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides \
0	0.142473	2.188833	-2.192833	-0.744778	0.569958
1	0.451036	3.282235	-2.192833	-0.597640	1.197975
2	0.451036	2.553300	-1.917553	-0.660699	1.026697
3	3.073817	-0.362438	1.661085	-0.744778	0.541412
4	0.142473	2.188833	-2.192833	-0.744778	0.569958

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates \
0	-1.100140	-1.446359	1.034993	1.813090	0.193097
1	-0.311320	-0.862469	0.701486	-0.115073	0.999579
2	-0.874763	-1.092486	0.768188	0.258120	0.797958
3	-0.762074	-0.986324	1.101694	-0.363868	0.327510
4	-1.100140	-1.446359	1.034993	1.813090	0.193097

```
alcohol
```


```
0 -0.915464
1 -0.580068
2 -0.580068
3 -0.580068
4 -0.915464
```

Determining the Elbow point using K means clustering

```
In [4]: sse = {}
```

```
for k in range(1, 20):
    kmeans = KMeans(n_clusters=k, max_iter=1000).fit(wine_data_scaled)
    wine_data["clusters"] = kmeans.labels_
    sse[k] = kmeans.inertia_
```

```
plt.figure()
plt.plot(list(sse.keys()), list(sse.values()))
plt.scatter(list(sse.keys()), list(sse.values()))
plt.title('elbow plot')
plt.xlabel("Number of clusters")
plt.ylabel("SSE")
plt.show()
```



From the above we can observe that $K=5/6$ is the perfect choice of K . Thus the plot shows that the number of groups to choose is 5. Hence let's run K means algorithm for $k=5$ and find clusters in the data.

2.3 Making Pair wise profiling plots and labelling wines with respect to its ingredients


plotting alcohol vs pH clusters

```
In [5]: wine_data.clusters.unique()
```


```
Out[5]: array([ 0, 18,  4, 16,  9,  7,  3, 11,  6, 15, 17,  8,  1,  5, 13,  2, 10,
                12, 14])
```

```
In [6]: plt.figure(figsize=(15,15))
plt.scatter(wine_data_scaled.pH,
            wine_data_scaled.alcohol,
            c=wine_data.clusters,
            cmap='inferno')

plt.legend(wine_data.clusters.unique())
plt.xlabel('pH')
plt.ylabel('alcohol')
plt.show()
```



```
In [7]: plt.figure(figsize=(15,15))
plt.scatter(wine_data_scaled.pH,
            wine_data_scaled.sulphates,
            c=wine_data.clusters,
            cmap='inferno')
plt.xlabel('pH')
plt.ylabel('sulphates')
plt.show()
```




```
In [8]: plt.figure(figsize=(15,15))
plt.scatter(wine_data_scaled.pH,
```

```

wine_data_scaled['total sulfur dioxide'],
c=wine_data.clusters,
cmap='inferno')
plt.xlabel('pH')
plt.ylabel('Total sulfur dioxide')
plt.show()


```




```

In [10]: plt.figure(figsize=(15,15))
plt.scatter(wine_data_scaled.alcohol,
            wine_data_scaled.sulphates,
            c=wine_data.clusters,
            cmap='inferno')
plt.xlabel('alcohol')
plt.ylabel('sulphates')
plt.show()


```



```
In [11]: plt.figure(figsize=(15,15))
plt.scatter(wine_data_scaled.alcohol,
            wine_data_scaled['total sulfur dioxide'],
            c=wine_data.clusters,
            cmap='inferno')
plt.xlabel('alcohol')
plt.ylabel('Total sulfur dioxide')
plt.show()
```



```
In [12]: plt.figure(figsize=(15,15))
plt.scatter(wine_data_scaled.sulphates,
            wine_data_scaled['total sulfur dioxide'],
            c=wine_data.clusters,
            cmap='inferno')
plt.xlabel('Sulfates')
plt.ylabel('Total Sulfur dioxide')
plt.show()
```

3 KNN Classifier

```
In [59]: red_wine_data = pd.read_csv('../Dataset/winequality-red.csv', sep=';')
white_wine_data = pd.read_csv('../Dataset/winequality-white.csv', sep=';')

wine_data = pd.concat([red_wine_data, white_wine_data])
bins = (2, 5, 6, 10)
group_names = ['bad', 'medium', 'good']
wine_data['quality'] = pd.cut(wine_data['quality'], bins = bins, labels = group_names)
wine_data.iloc[:, :11].describe()
```

```
Out [59]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	\
count	6497.000000	6497.000000	6497.000000	6497.000000	

mean	7.215307	0.339666	0.318633	5.443235
std	1.296434	0.164636	0.145318	4.757804
min	3.800000	0.080000	0.000000	0.600000
25%	6.400000	0.230000	0.250000	1.800000
50%	7.000000	0.290000	0.310000	3.000000
75%	7.700000	0.400000	0.390000	8.100000
max	15.900000	1.580000	1.660000	65.800000

	chlorides	free sulfur dioxide	total sulfur dioxide	density \
count	6497.000000	6497.000000	6497.000000	6497.000000
mean	0.056034	30.525319	115.744574	0.994697
std	0.035034	17.749400	56.521855	0.002999
min	0.009000	1.000000	6.000000	0.987110
25%	0.038000	17.000000	77.000000	0.992340
50%	0.047000	29.000000	118.000000	0.994890
75%	0.065000	41.000000	156.000000	0.996990
max	0.611000	289.000000	440.000000	1.038980

	pH	sulphates	alcohol
count	6497.000000	6497.000000	6497.000000
mean	3.218501	0.531268	10.491801
std	0.160787	0.148806	1.192712
min	2.720000	0.220000	8.000000
25%	3.110000	0.430000	9.500000
50%	3.210000	0.510000	10.300000
75%	3.320000	0.600000	11.300000
max	4.010000	2.000000	14.900000

```
In [60]: scaler = StandardScaler()
```

```
data = wine_data.iloc[:,11].values
scaled_features = scaler.fit_transform(data)
```

```
wine_data_scaled = pd.DataFrame(scaled_features, index=wine_data.index, columns=wine_data.columns)
wine_data_scaled.head()
```

```
Out[60]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides \
0	0.142473	2.188833	-2.192833	-0.744778	0.569958
1	0.451036	3.282235	-2.192833	-0.597640	1.197975
2	0.451036	2.553300	-1.917553	-0.660699	1.026697
3	3.073817	-0.362438	1.661085	-0.744778	0.541412
4	0.142473	2.188833	-2.192833	-0.744778	0.569958

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates \
0	-1.100140	-1.446359	1.034993	1.813090	0.193097
1	-0.311320	-0.862469	0.701486	-0.115073	0.999579
2	-0.874763	-1.092486	0.768188	0.258120	0.797958
3	-0.762074	-0.986324	1.101694	-0.363868	0.327510

```
4          -1.100140          -1.446359  1.034993  1.813090  0.193097
```

```
    alcohol
0 -0.915464
1 -0.580068
2 -0.580068
3 -0.580068
4 -0.915464
```

```
In [83]: X_train, X_test, Y_train, Y_test = train_test_split(wine_data.iloc[:, :11],
                                                             wine_data.quality,
                                                             test_size=0.2,
                                                             random_state=42)
```

```
In [84]: knn = KNeighborsClassifier(20)
         knn.fit(X_train, Y_train)
```

```
Out[84]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                              metric_params=None, n_jobs=1, n_neighbors=20, p=2,
                              weights='uniform')
```

```
In [85]: knn.score(X_train, Y_train)
```

```
Out[85]: 0.8166249759476621
```

```
In [89]: knn_score = []
```

```
for i in range(1,160):
    knn = KNeighborsClassifier(i)
    knn.fit(X_train, Y_train)
    knn_score.append(knn.score(X_train, Y_train))

plt.plot(range(1, 160), knn_score)
plt.xlabel("KNN score")
plt.ylabel("Number of neighbours")
plt.show()
```

